

# Concluding Exercise in R

Sharon Hadar

09.13.2017

## Exercise #1:

In this exercise you will perform a Monte Carlo study of a simple linear regression. Let  $\beta_0 = 2$  and  $\beta_1 = 1$  and assume the following relationship holds:  $y = \beta_0 + \beta_1 x + \text{epsilon}$ , where  $x \in [0, 10]$  and  $\text{epsilon} \in N(0, \sigma^2)$  with  $\sigma^2 = 3$ . For each combination of  $n$  and  $m$  where  $n \in \{20, 100\}$  and  $m \in \{10, 500\}$ , do the following:

Create a vector of  $n$  evenly spaced values for  $x$  in the range  $[0, 10]$ , then generate  $n$  random observations of  $y$  according to the model above

```
Xvector <- function(n)
{
  return (seq(0, 10, length.out = n))
}

Xvector(5)
```

```
## [1] 0.0 2.5 5.0 7.5 10.0
```

```
Yvector <- function(x)
{
  epsilon = rnorm(length(x), mean = 0, sd = sqrt(3))
  y = 2 + 1*x + epsilon
  return (y)
}

Yvector(Xvector(5))
```

```
## [1] 2.901668 6.362925 7.181206 10.185965 9.953682
```

Estimate the values of  $\beta_0$  and  $\beta_1$  using the least squares method (lm function).

```
CreateDataFrame <- function(n)
{
  df <- data.frame(Xvector(n), Yvector(Xvector(n)))
  names(df) <- c("X", "Y")
  return (df)
}

GetLm <- function(n)
{
  return (model = lm(data = CreateDataFrame(n), formula = Y~X))
}

GetLm(5)$coefficients
```

```
## (Intercept)          X
##      2.445437      0.935845
```

Repeat steps a) and b) m times.

```
GetCoefficientsDf <- function(m,n)
{
  if(m<1){
    return (NULL)
  }
  df <- data.frame(X= numeric(0), Y= numeric(0))
  for(model in 1:m)
  {
    de <- GetLm(n)$coefficients
    df <- rbind(df, de)
  }
  names(df)<-c("β0","β1")
  return (df)
}

GetCoefficientsDf(2,12)
```

```
##      <U+03B2>0 <U+03B2>1
## 1  1.1098263  1.0036383
## 2  0.6133471  0.9708139
```

Calculate the mean and variance of the estimates for  $\beta_0$  and  $\beta_1$  over the m simulations.

```
GetMeanSD <-function(m,n)
{
  df = GetCoefficientsDf(m,n)
  sd = apply(df, 2, sd)
  mean = apply(df, 2, mean)
  var = sd*sd
  return(rbind(mean,var))
}

GetMeanSD(100,100)
```

```
##      <U+03B2>0      <U+03B2>1
## mean 1.9852021 0.997909728
## var  0.1012728 0.003375533
```

Create a table summarising the results that includes the columns: n, m, beta0\_mean, beta1\_mean, beta0\_var and beta1\_var. Explain your findings regarding the effect of the size of n and the size of m on the results.

```
CreateSummarisingTable <- function(m,n)
{
  exData = GetMeanSD(m,n)
  df =
data.frame(n,m,exData["mean","β0"],exData["mean","β1"],exData["var","β0"],exData["var","β1"])
  names(df)<-c("n", "m", "beta0_mean","beta1_mean", "beta0_var", "beta1_var")
  return(df)
}

CreateSummarisingTable(100,100)
```

```
##      n   m beta0_mean beta1_mean beta0_var  beta1_var
## 1 100 100   1.993026   1.004109 0.1308815 0.00427993
```

the best estimation would be  $\beta_0=2$  &  $\beta_1=1$ . we can see that the correlation is positive-> for  $n \rightarrow \infty$   $m \rightarrow \infty$  we would get  $\beta_0=2$  &  $\beta_1=1$ . that is because in both cases epsilon would be 0 and the simulation would have the same formula as the linear model.

```
library(ggplot2)
```


```
## Warning: package 'ggplot2' was built under R version 3.4.1
```

```
library(gridExtra)
```

```
## Warning: package 'gridExtra' was built under R version 3.4.1
```

```
seeResult <- function(Maxm,Maxn)
{
  df <- data.frame(n= Maxm,m= Maxm,beta0_mean= 0,beta1_mean= 0, beta0_var= 0, beta1_var= 0)
  for(m in 3:Maxm){
    for(n in 3:Maxn)
    {
      addone = CreateSummarisingTable(m,n)
      addone$beta0_mean = addone$beta0_mean-2
      addone$beta1_mean = addone$beta1_mean-1
      df=rbind(df,addone)
    }
  }
  return (df)
}

df<-seeResult(50,50)
beta0_mean<-ggplot(data =df)+
  geom_raster(mapping = aes(n,m,fill= abs(beta0_mean)), data = df)
beta1_mean<-ggplot(data =df)+
  geom_raster(mapping = aes(n,m,fill= abs(beta1_mean)), data = df)
beta0_var<-ggplot(data =df)+
  geom_raster(mapping = aes(n,m,fill= abs(beta0_var)), data = df)
beta1_var<-ggplot(data =df)+
  geom_raster(mapping = aes(n,m,fill= abs(beta1_var)), data = df)
grid.arrange(beta0_mean, beta1_mean,beta0_var,beta1_var)
```



## answer

we can see in the figure 4 models for each estimation value with incrementing n and m.

in mean estimators the value in each pixel represented as the absolute distance from the real mean. where in sd estimators it is represented as the same value.

from this figure we can see the gradient of change in value: in all 4 cases the gradient direction is very similar to the 'n' axis, meaning 'n' margin is better then 'm',

we can also see that the mean estimators gradient direction gets more influence from 'm' meaning that the mean estimation relay more on the number of iteration averaged than the sd estimator.

## Exercise #2:

In this exercise you will implement the K-nearest-neighbor (KNN) method for regression. Let  $d(a, b) = \sqrt{\sum_{i=1}^p (a_i - b_i)^2}$  be the Euclidean distance between points  $a$  and  $b$ . Assume we have a training set of  $x_1, \dots, x_n$  points (observations with  $p \geq 1$  dimensions) and outcome values  $y_1, \dots, y_n$ , and we are given  $x_0$  as a new test point. Define  $d_i(x_0) = d(x_0, x_i)$  as the distance between  $x_0$  and each point  $x_i$ . Let  $d^{(i)}(x_0)$  be the distance between  $x_0$  and the  $i$ 'th nearest point to  $x_0$  (i.e.,  $d^{(1)}(x_0)$  is the distance between  $x_0$  and the nearest point,  $d^{(2)}(x_0)$  is the distance between  $x_0$  and the second nearest point. . . ). The set of K-nearest neighbors of  $x_0$  is given by  $N_k(x_0) = \{x_i: d_i(x_0) \leq d^{(K)}(x_0)\}$ . The KNN estimator of  $y_0$  is given by  $\hat{y}_0 = \frac{1}{K} \sum_{i \in N_k(x_0)} y_i$ .

Create a function in R that calculates the KNN estimator at a point  $x_0$ . The function should get as input the points (matrix of size  $n \times p$ ), the outcomes ( $n \times 1$ ), the new point  $x_0$  ( $1 \times p$ ) and the number K. The output should be the estimate for  $y_0$ .

```

calcDistance <- function(a,b)
{
  if(length(a)!=length(b))
  {
    return (NULL)
  }
  vac = (a-b)^2
  value <- sum(vac)
  return (sqrt(value))
}

GetDistanceMatrix <- function(points,newpoint,outcomes)
{
  numberOfRow = nrow(points)
  returnValue = matrix(0, nrow = numberOfRow, ncol = 1)
  for(index in 1:numberOfRow)
  {
    ans = calcDistance(points[index, ],newpoint)
    returnValue[index,1] =ans
    if(is.null(ans))
    {
      print("error while calc distance")
    }
  }
  return (data.frame(returnValue,outcomes))
}

KNNestimator <- function(points,outcomes,newpoint,knumber)
{
  data <- GetDistanceMatrix(points,newpoint,outcomes)
  head <- head(data[order(data[,1]),], knumber)

  if(is.character(outcomes))
  {
    slice(
      arrange(
        count(head, vars = outcomes)
        ,desc(n))
      , 1 ) -> firstk
    return (firstk[1,1]$vars)
  }
  else if(is.double(outcomes))
  {
    return(sum(head$outcomes)/knumber)
  }
  else
  {
    return("Dont support outcome type")
  }
}

sizeFactor = 50
k=50
propvec = c(0.25, 0.25, 0.25, 0.25)
points=matrix(rexp(sizeFactor*sizeFactor), sizeFactor)

```

```

newpoint = rexp(sizeFactor)
characterOutcomes = sample( LETTERS[1:4], sizeFactor, replace=TRUE, prob=propvec)
doubleOutcomes = rexp(sizeFactor)

(KNNestimator(points,characterOutcomes,newpoint,k))

```

```

## [1] A
## Levels: A B C D

```

```

(KNNestimator(points,doubleOutcomes,newpoint,k))

```

```

## [1] 0.8314869

```

Generate  $x$  and  $y$  using the setup from question 1 with  $n = 100$ . Calculate the KNN estimator with  $K = 1, 5, 9$  for the points  $x_0 = x_1, \dots, x_n$  (i.e., for each point  $x_i$ , estimate  $y_i$  using all the points in the training set). For each  $K$ , create a ggplot with the true function ( $y_i$  without the noise term - drawn using a black line), the observations ( $y_i$  - drawn using black circles) and the KNN estimators ( $\hat{y}_i$  - drawn using a red line).

```

Xvector <- function(n)
{
  return (seq(0, 10, length.out = n))
}

Yvector <- function(x)
{
  set.seed(123)
  epsilon = rnorm(n=length(x), mean = 0, sd = sqrt(3))
  y = 2 + 1*x + epsilon
  return (y)
}

n=100

xdata = as.data.frame(Xvector(n))
ydata = Yvector(Xvector(n))
ytrend = 2 + 1*Xvector(n)


returnValue <- data.frame(x= xdata, y= ydata,
ytrend,est_y_k1=numeric(n),est_y_k5=numeric(n),est_y_k9=numeric(n))
for(row in 1:nrow(returnValue))
{
  x0<-  returnValue[row,'Xvector.n.']
  returnValue[row,'est_y_k1']<-KNNestimator(xdata,ydata,x0,1)
  returnValue[row,'est_y_k5']<-KNNestimator(xdata,ydata,x0,5)
  returnValue[row,'est_y_k9']<-KNNestimator(xdata,ydata,x0,9)
}
colnames(returnValue) <- c("X", "Y", "Trend", "k1", "k5", "k9")
head(returnValue)

```


##	X	Y	Trend	k1	k5	k9
## 1	0.0000000	1.029228	2.000000	1.029228	2.537294	2.633425
## 2	0.1010101	1.702331	2.101010	1.702331	2.537294	2.633425
## 3	0.2020202	4.901782	2.202020	4.901782	2.537294	2.633425
## 4	0.3030303	2.425154	2.303030	2.425154	3.426574	2.633425
## 5	0.4040404	2.627973	2.404040	2.627973	3.766986	2.633425
## 6	0.5050505	5.475630	2.505051	5.475630	2.889814	2.756531

```
gatheredValues = gather(data = returnValue, key = K, value = estimation, ... = k1:k9, factor_key = TRUE)
```

```
ggplot(data = gatheredValues) +
  geom_point(mapping = aes(X, estimation)) +
  geom_smooth(mapping = aes(x = gatheredValues$X, y = gatheredValues$Trend), method = 'lm') +
  facet_wrap(~K)
```



```
ggplot(data = gatheredValues) +
  geom_point(mapping = aes(X, Y)) +
  geom_smooth(mapping = aes(x = gatheredValues$X, y = gatheredValues$Trend), method = 'lm') +
  geom_smooth(mapping = aes(x = gatheredValues$X, y = gatheredValues$estimation),
    method = 'lm', color = 'red') +
  facet_wrap(~K)
```



c. Explain how would you go about selecting a value for K to use for prediction using the training set.

->Select the k that minimize the error can use cross validation to get the optimal value.

## Exeercise #3:

3. In this exercise you will employ logistic regression with the '2016 World Health Report' data set, to examine how well can we predict if a country is European or African based on the given predictors.
  - a. Using the 'dplyr' package or any other programming method you want to apply, create a subset of the '2016 World Health Report' data set, that includes only European or African countries (50 European countries and 43 African countries). The new data set will have a new field called 'Region2' that will contain either 'Europe' or 'Africa', as well as a field called 'IsEuropean' which is set to true for European countries and false otherwise.

```
world <- read.csv('./2016 World Happiness Report.csv',header=T)
attach(world)
summary(world$Region)
```



```
##      Australia and New Zealand      Central and Eastern Europe
##              2                      29
##              Eastern Asia      Latin America and Caribbean
##              6                      24
## Middle East and Northern Africa      North America
##              19                      2
##              Southeastern Asia      Southern Asia
##              9                      7
##              Sub-Saharan Africa      Western Europe
##              38                      21
```

```
EuroGroup =c("Central and Eastern Europe","Western Europe")
AfroGroup =c("Middle East and Northern Africa","Sub-Saharan Africa")

df2 <- filter(.data = world,Region %in% c(EuroGroup, AfroGroup) ) %>%
  mutate(Region2=if_else(Region %in% EuroGroup, "Europe", "Africa", missing = NULL),
         IsEuropean=if_else(Region %in% EuroGroup, TRUE, FALSE, missing = NULL))
```

```
## Warning: package 'bindrcpp' was built under R version 3.4.1
```

```
df2$Region2 <- as.factor(df2$Region2)

summary(df2)
```

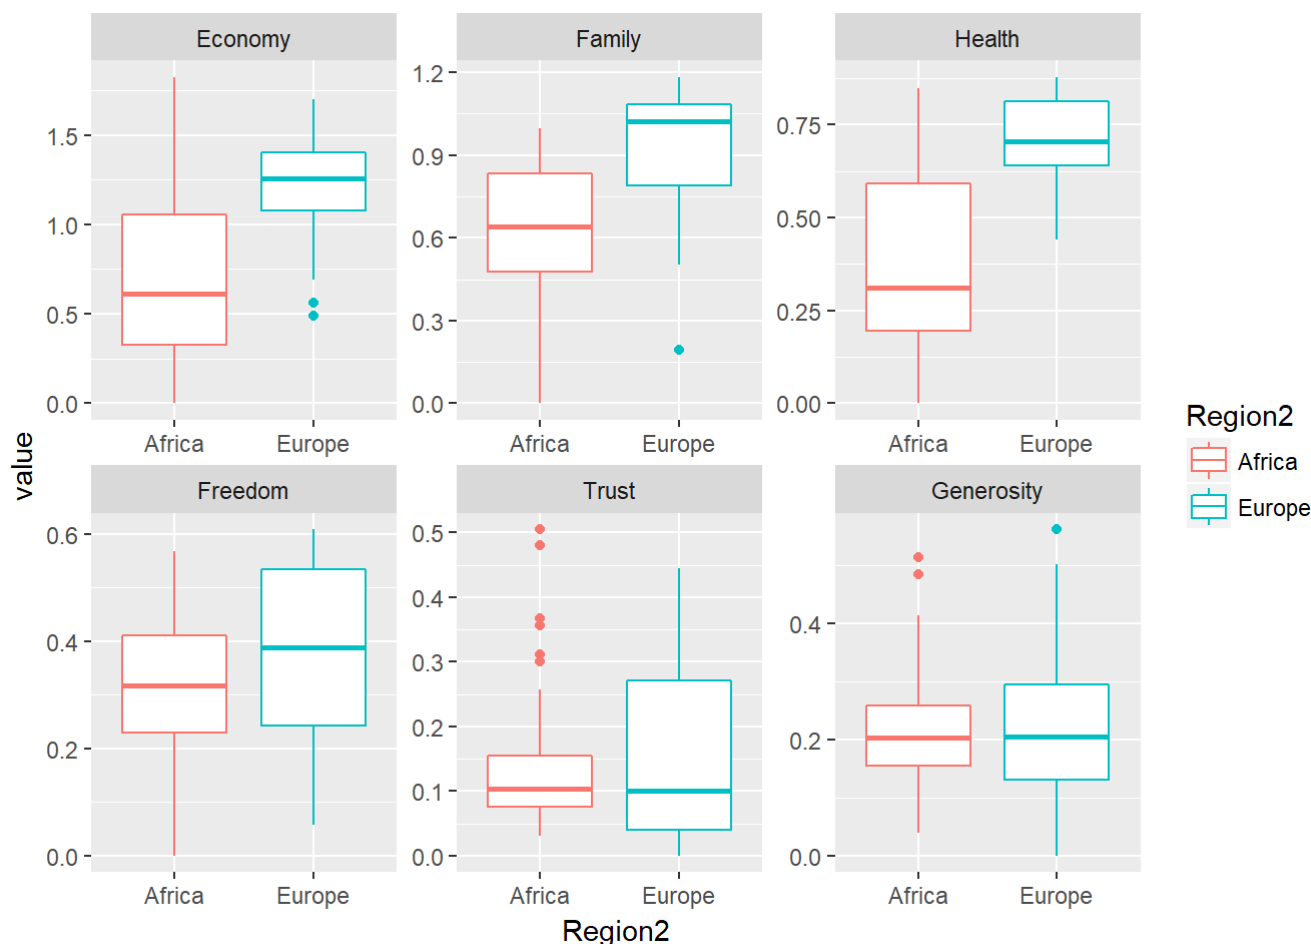
```
##      Country      Region      Happiness
## Albania   : 1 Sub-Saharan Africa :38 Min. :2.905
## Algeria   : 1 Central and Eastern Europe :29 1st Qu.:4.228
## Angola    : 1 Western Europe :21 Median :5.145
## Armenia   : 1 Middle East and Northern Africa:19 Mean :5.193
## Austria   : 1 Australia and New Zealand : 0 3rd Qu.:5.948
## Azerbaijan: 1 Eastern Asia : 0 Max. :7.526
## (Other)   :101 (Other) : 0
## Economy      Family      Health      Freedom
## Min. :0.0000 Min. :0.0000 Min. :0.0000 Min. :0.0000
## 1st Qu.:0.5578 1st Qu.:0.6251 1st Qu.:0.3006 1st Qu.:0.2347
## Median :1.0344 Median :0.8068 Median :0.5947 Median :0.3602
## Mean :0.9328 Mean :0.7725 Mean :0.5279 Mean :0.3422
## 3rd Qu.:1.2896 3rd Qu.:1.0065 3rd Qu.:0.7236 3rd Qu.:0.4387
## Max. :1.8243 Max. :1.1833 Max. :0.8790 Max. :0.6085
##
## Trust      Generosity      Region2 IsEuropean
## Min. :0.0000 Min. :0.0000 Africa:57 Mode :logical
## 1st Qu.:0.0540 1st Qu.:0.1389 Europe:50 FALSE:57
## Median :0.1039 Median :0.2036 TRUE :50
## Mean :0.1409 Mean :0.2197
## 3rd Qu.:0.1786 3rd Qu.:0.2774
## Max. :0.5052 Max. :0.5624
##
```

is more than 43 since it is mixed with middle east countries but i assume and hope this is not the point of this exercise

- b. Using 'ggplot2' and with the help of 'tidyr', create a plot similar to Figure1 below. Based on this figure, which predictors could be helpful in predicting if a country is from Europe or Africa?

```
world_long = gather(data = df2, key = sector, value = value, ... = Economy:Generosity, factor_key = TRUE)
```

```
ggplot(data = world_long) +  
  geom_boxplot(mapping = aes(x = Region2, y = value, color = Region2)) +  
  facet_wrap(~sector, scales = "free")
```



best properties base on this plot are health, economy, and family.

c) Using glm, fit a simple logistic regression to `IsEuropean` using each of the six predictors shown in Figure 1. According to the fits, which predictor by itself is better in distinguishing between European and African countries? Create a ggplot showing the probability of a country being European according to the value of this best predictor.

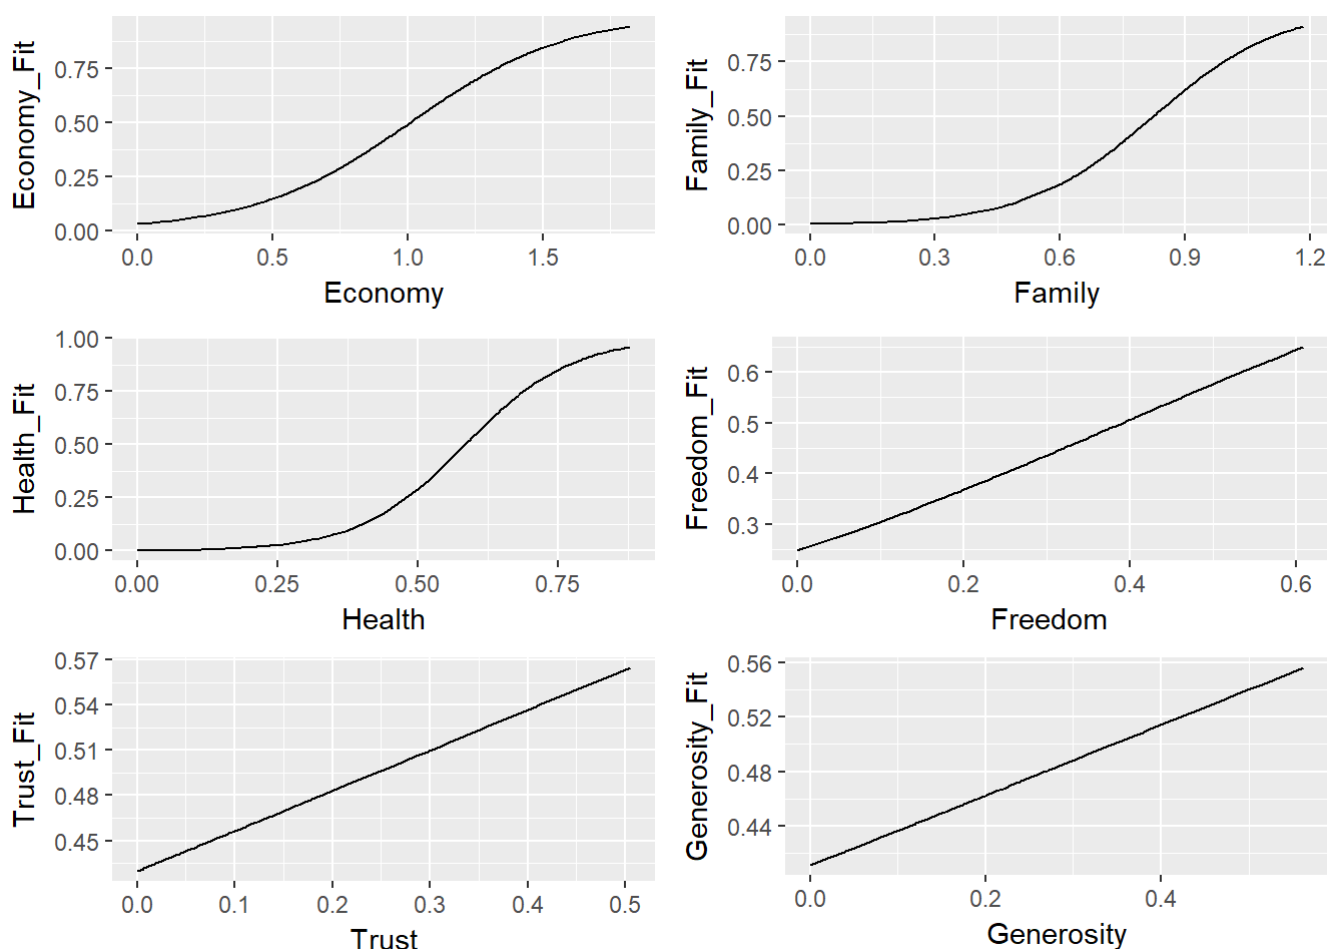
```
library(ggplot2)  
attach(df2)
```

```
## The following objects are masked from world:  
##  
## Country, Economy, Family, Freedom, Generosity, Happiness,  
## Health, Region, Trust
```

```

logm.Economy<-glm(formula = IsEuropean ~ Economy, family =binomial , data = df2)
logm.Family<-glm(formula = IsEuropean ~ Family, family =binomial , data = df2)
logm.Health<-glm(formula = IsEuropean ~ Health, family =binomial , data = df2)
logm.Freedom<-glm(formula = IsEuropean ~ Freedom, family =binomial , data = df2)
logm.Trust<-glm(formula = IsEuropean ~ Trust, family =binomial , data = df2)
logm.Generosity<-glm(formula = IsEuropean ~ Generosity, family =binomial , data = df2)
Economy_Fit=fitted(logm.Economy)
Family_Fit=fitted(logm.Family)
Health_Fit=fitted(logm.Health)
Freedom_Fit=fitted(logm.Freedom)
Trust_Fit=fitted(logm.Trust)
Generosity_Fit=fitted(logm.Generosity)
g1 <-ggplot() + geom_line(mapping=aes(x=Economy,y=Economy_Fit))
g2 <-ggplot() + geom_line(mapping=aes(x=Family,y=Family_Fit))
g3 <-ggplot() + geom_line(mapping=aes(x=Health,y=Health_Fit))
g4 <-ggplot() + geom_line(mapping=aes(x=Freedom,y=Freedom_Fit))
g5 <-ggplot() + geom_line(mapping=aes(x=Trust,y=Trust_Fit))
g6 <-ggplot() + geom_line(mapping=aes(x=Generosity,y=Generosity_Fit))
grid.arrange(g1,g2,g3,g4,g5,g6)

```



we can see that health hold the best separator between 0 probability to be European and 1 - it is the most similar to sigmoid function d) Based on the fit using the best predictor from c), create a confusion table of actual region vs. predicted region, using each of the following threshold values for  $p(\text{IsEuropean}=\text{Yes})$ : 1/4, 1/2, 3/4. Report the sensitivity and specificity of the classifier in each case. Which of these classifiers would you use? Specify if it depends on anything.

```

confusion.glm <- function(data, model,tresh) {
  prediction <- ifelse(predict(model, data, type='response') > tresh, TRUE, FALSE)
  confusion <- table(prediction, as.logical(model$y))
  confusion <- cbind(confusion, c(1 - confusion[1,1]/(confusion[1,1]+confusion[2,1]), 1 - co
nfusion[2,2]/(confusion[2,2]+confusion[1,2])))
  confusion <- as.data.frame(confusion)
  names(confusion) <- c('FALSE', 'TRUE', 'class.error')
  confusion
}
confusion.glm(df2,logm.Health,0.25)

```

```

##          FALSE TRUE class.error
## FALSE      38    1  0.3333333
## TRUE       19   49  0.0200000

```

```

confusion.glm(df2,logm.Health,0.5)

```

```

##          FALSE TRUE class.error
## FALSE      42    8  0.2631579
## TRUE       15   42  0.1600000

```

```

confusion.glm(df2,logm.Health,0.75)

```

```

##          FALSE TRUE class.error
## FALSE      53   22  0.07017544
## TRUE        4   28  0.44000000

```

## Exercise #4:

4. In this exercise you will employ the 'caret' package wrapper for randomForest with the '2016 World Health Report' data set, in order to fit and tune a model for predicting the Happiness score of a country.
  - a. Use `caret::createDataPartition` to create training and test sets, with the training set containing 70% of the data.

```

world <- read.csv('./2016 World Happiness Report.csv',header=T)
attach(world)

```

```

## The following objects are masked from df2:
##
## Country, Economy, Family, Freedom, Generosity, Happiness,
## Health, Region, Trust

```

```

## The following objects are masked from world (pos = 5):
##
## Country, Economy, Family, Freedom, Generosity, Happiness,
## Health, Region, Trust

```

```
#install.packages("caret")
#install.packages("caret", dependencies = c("Depends", "Suggests"))

library("caret")
```

```
## Warning: package 'caret' was built under R version 3.4.1
```

```
## Loading required package: lattice
```

```
train<-createDataPartition(y=Happiness, p = 0.7, list = FALSE)
world.train <- world[train,]
world.test <- world[-train,]
Happiness.test <- world$Happiness[-train]
```

- b. Explain what is the meaning of the tuning parameter 'mtry' of randomForest. What is the maximum number it can take in our case, assuming we want to fit Happiness using all the predictors in the data set (not including the field 'Country')?

mtry is the number of variables randomly sampled as candidates at each split, the maximum number it can take is bound by the number of variables in the model, as it specifies the size of the variable subset that is randomly picked for each random forest iteration. the maximum number would be all distinct factors in each factor feature plus the number of numeric features:

```
numericValues = table(sapply(world, is.numeric))["TRUE"]
factorValues = sapply(world[,sapply(world, is.factor)], nlevels)
(maxmtry = sum(factorValues,numericValues))
```

```
## [1] 174
```

maximum value for mtry =174

- c. Use caret::train with method='rf' to fit the model to the training set. Set that the method will search for the best possible value for 'mtry' out of all possible values, using root-mean-square error (RMSE) as the metric for selecting the best value. Set the resampling method to be used to 10-fold cross-validation. Set the number of trees to be grown to 500. Make sure that the importance of the predictors will be assessed. Using the obtained fit, plot the cross-validation RMSE as a function of mtry and the cross-validation Rsquared as a function of mtry.

What is the selected value for 'mtry' according to the fit? What would have been the selected value for 'mtry' if we used Rsquared as the metric for selection?

```
#install.packages("randomForest")
#library(plyr)
set.seed(123)
Happiness.rfFit <- train(Happiness~.,world.train,
                        method='rf',
                        metric='RMSE',
                        trControl=trainControl(method='cv',number=10),
                        tuneGrid=data.frame(mtry=3:150),
                        ntree=500,
                        importance=TRUE)
```

```
## Warning: package 'randomForest' was built under R version 3.4.1
```

```
## randomForest 4.6-12
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##  
## Attaching package: 'randomForest'
```

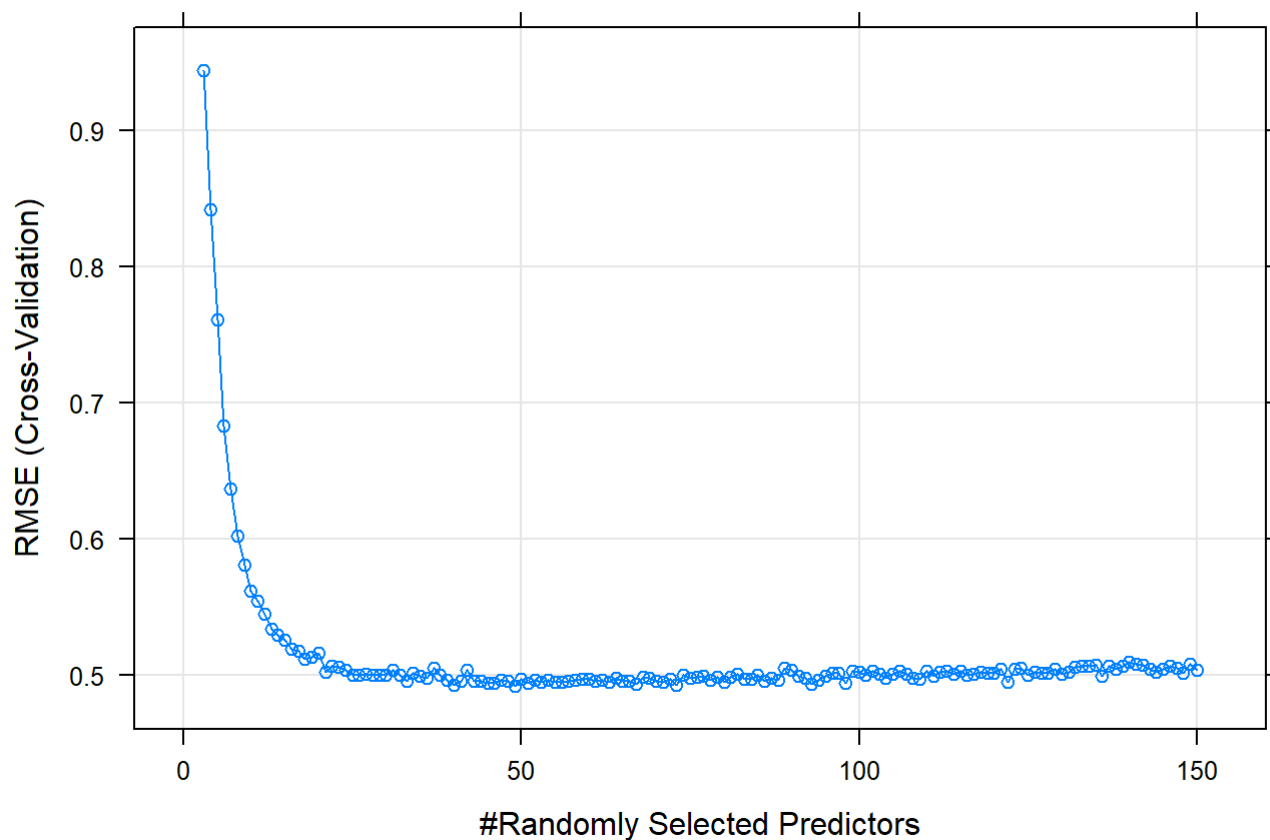
```
## The following object is masked from 'package:dplyr':  
##  
## combine
```

```
## The following object is masked from 'package:gridExtra':  
##  
## combine
```

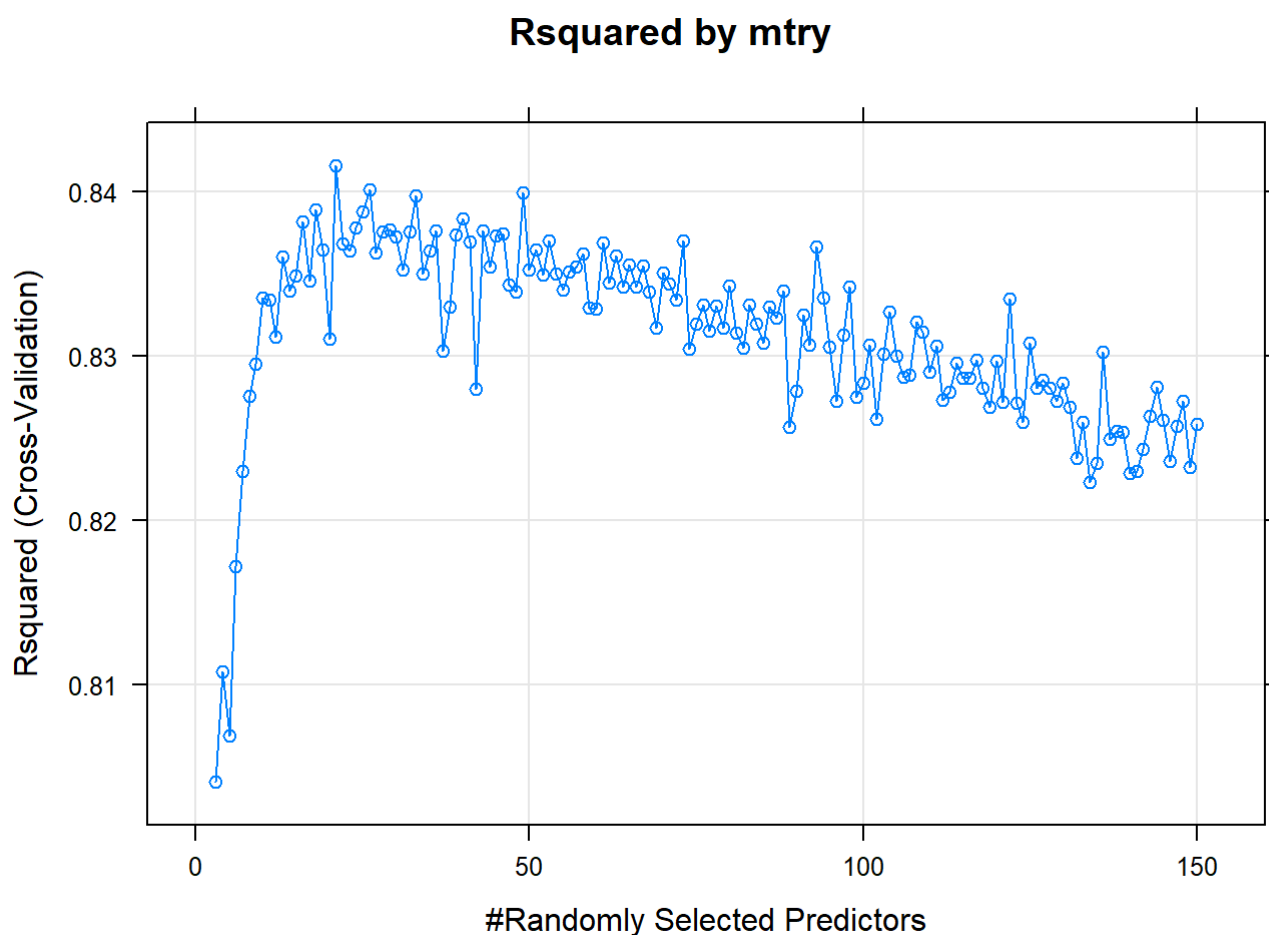
```
## The following object is masked from 'package:ggplot2':  
##  
## margin
```

```
plot(Happiness.rfFit, metric='RMSE',main="RMSE by mtry")
```

## RMSE by mtry



```
plot(Happiness.rfFit, metric='Rsquared',main="Rsquared by mtry")
```



```
getTrainPerf(Happiness.rfFit)
```

```
##   TrainRMSE TrainRsquared TrainMAE method
## 1 0.4918707    0.8399018 0.3980317    rf
```

```
(winning.RMSE = Happiness.rfFit$results[which.min(Happiness.rfFit$results[, "RMSE"]), ])
```

```
##   mtry    RMSE Rsquared    MAE   RMSESD RsquaredSD    MAESD
## 47   49 0.4918707 0.8399018 0.3980317 0.09740359 0.05667068 0.07089441
```

```
(winning.Rsquared = Happiness.rfFit$results[which.max(Happiness.rfFit$results[, "Rsquared"]),
])
```

```
##   mtry    RMSE Rsquared    MAE   RMSESD RsquaredSD    MAESD
## 19   21 0.5021744 0.8415588 0.4081469 0.09216605 0.05829391 0.06812168
```

```
mtryw = winning.Rsquared$mtry
```

the wining mtry for Rsquared is mtry= 21


- d. Using the obtained fit, plot the calculated variable importance. What are the top-3 predictors effecting Happiness according to this plot?

```

ImpMeasure<-data.frame(varImp(Happiness.rfFit)$importance)
ImpMeasure$Vars<-row.names(ImpMeasure)
ImpMeasure<-ImpMeasure[order(-ImpMeasure$Overall),][1:20,]

barplot(ImpMeasure$Overall, names.arg=ImpMeasure$Vars ,main="important Measure by variable",las=2)

```



```
top = head(ImpMeasure,3)$Vars
```

the top-3 are :Economy, Health, Family

- e. Use the model fit to estimate the Happiness score for the test set. Calculate the RMSE for these predictions. Create a ggplot with the actual Happiness score vs. the predicted Happiness score.

```

Happiness.rf <- predict(Happiness.rfFit, newdata=world.test)
(test.mse.rf <- mean((Happiness.rf - Happiness.test)^2))

```

```
## [1] 0.3144871
```

```

range = (min(Happiness.test)-0.25):(max(Happiness.test)+0.75)
ggplot()+
  geom_point(mapping = aes(x=Happiness.test,y = Happiness.rf))+
  ylab("Predicted Happiness")+
  xlab("Actual Happiness")+
  geom_smooth(mapping = aes(x = range, y = range), method='lm')+
  geom_segment(aes(x=Happiness.test, xend=Happiness.test, y=Happiness.rf,
yend=Happiness.test, color="error"))

```



